

February 20, 2025

```
[1]: #
import pandas as pd
import numpy as np

# xlsx
# 'your_file.xlsx'
df = pd.read_excel(r"D:\\\\Data2017 -2019_1.xlsx")

#
print(df.head())

#
print(df.info())
```

	X	Y	produkt	date	Agrotech	Soil	Sum	Ca	Mg	\
0	59.423807	30.038469	0.1	2018	K	2.31	8.38	7.12	1.26	
1	59.423821	30.038704	0.1	2018	K	2.31	8.38	7.12	1.26	
2	59.423848	30.039059	0.1	2018	K	2.31	8.38	7.12	1.26	
3	59.423862	30.039295	0.1	2018	K	2.31	8.38	7.12	1.26	
4	59.423868	30.039537	0.1	2018	K	2.31	8.38	7.12	1.26	

	pH_KCL	P	K	N	Org
0	5.4	327	188	0.31	0.89
1	5.4	327	188	0.31	0.89
2	5.4	327	188	0.31	0.89
3	5.4	327	188	0.31	0.89
4	5.4	327	188	0.31	0.89

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 8182 entries, 0 to 8181

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	X	8182 non-null	float64
1	Y	8182 non-null	float64
2	produkt	8182 non-null	float64
3	date	8182 non-null	int64
4	Agrotech	8182 non-null	object
5	Soil	8171 non-null	object

```

6          8182 non-null   float64
7   Sum      8182 non-null   float64
8   Ca       8182 non-null   float64
9   Mg       8182 non-null   float64
10  pH_KCL   8182 non-null   float64
11  P        8182 non-null   int64
12  K        8182 non-null   int64
13  N        8182 non-null   float64
14  Org      8182 non-null   float64
dtypes: float64(10), int64(3), object(2)
memory usage: 959.0+ KB
None

```

```

[2]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.metrics import mean_squared_error, r2_score
from catboost import CatBoostRegressor
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats

# MAPE
def calculate_mape(y_true, y_pred):
    y_true, y_pred = np.array(y_true), np.array(y_pred)
    return np.mean(np.abs((y_true - y_pred) / y_true)) * 100

# 1.
def analyze_outliers(df, numeric_columns, z_score_threshold=3):
    outliers_info = {}
    all_outliers_indices = set()

    print("\n      :")
    for col in numeric_columns:
        z_scores = np.abs(stats.zscore(df[col]))
        outliers_mask = z_scores > z_score_threshold
        outliers_count = outliers_mask.sum()
        outliers_indices = df[outliers_mask].index

        outliers_info[col] = {
            'count': outliers_count,
            'percentage': (outliers_count / len(df)) * 100,
            'indices': outliers_indices
        }

    all_outliers_indices.update(outliers_indices)

```

```

        print(f"{col}: {outliers_count} ({outliers_info[col]['percentage']:.
↪2f}%)" )

    print(f"\n      : {len(all_outliers_indices)}")
    print(f"      : {(len(all_outliers_indices) / len(df)) * 100:.2f}%")

    return list(all_outliers_indices), outliers_info

# 2.
df['Soil'] = df['Soil'].fillna('Unknown')

#
numeric_features = ['', 'Sum', 'Ca', 'Mg', 'pH_KCL', 'P', 'K', 'N', 'Org',
↪'produkt']
#
print("      :", (df['produkt'] == 0).sum())

#
df = df[df['produkt'] > 0]
print("      :", len(df))
#
print("\n      :")
print(df['produkt'].describe())

df = df[df['produkt'] > 0]

print("\n      :")
print(df['produkt'].describe())

#
print("\n      :", len(df))
outliers_indices, outliers_info = analyze_outliers(df, numeric_features,
↪z_score_threshold=3)

#
remove_outliers = True

if remove_outliers:
    df_clean = df.drop(outliers_indices)
    print(f"      : {len(df_clean)}")
    print(f"      : {len(df) - len(df_clean)}")
else:
    df_clean = df
    print("      ")

# 3.
shuffle_index = np.random.permutation(len(df))

```

```

df_shuffled = df.iloc[shuffle_index].reset_index(drop=True)

features = ['Agrotech', 'Soil', '', 'Sum', 'Ca', 'Mg', 'pH_KCL', 'P', 'K', 'N',
            ↪ 'Org']
X = df_clean[features]
y = df_clean['produkt']

#
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
            ↪ random_state=42)

# 4. GridSearchCV
param_grid = {
    'learning_rate': [0.05, 0.1],
    'depth': [4, 6],
    'iterations': [500]
}

cat_features = ['Agrotech', 'Soil']

# GridSearchCV
base_model = CatBoostRegressor(
    cat_features=cat_features,
    random_seed=42,
    iterations=1000,
    depth=6,
    learning_rate=0.03,
    verbose=False
)

# GridSearchCV
grid_search = GridSearchCV(
    estimator=base_model,
    param_grid=param_grid,
    cv=3,
    scoring='neg_root_mean_squared_error',
    n_jobs=-1
)

print("\n ...")
#
grid_search.fit(X_train, y_train)

#
print("\n :")
print(grid_search.best_params_)

```

```

#
best_model = grid_search.best_estimator_
y_pred = best_model.predict(X_test)

# 5.
def print_metrics(y_true, y_pred, model_name):
    mse = mean_squared_error(y_true, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_true, y_pred)
    mape = calculate_mape(y_true, y_pred)

    print(f'\n {model_name}:')
    print(f'RMSE: {rmse:.4f}')
    print(f'R2 Score: {r2:.4f}')
    print(f'MAPE: {mape:.2f}%')

print_metrics(y_test, y_pred, "CatBoost ( )")

# 6.
feature_importance = pd.DataFrame({
    'feature': features,
    'importance': best_model.feature_importances_
})
print("\n :")
print(feature_importance.sort_values(by='importance', ascending=False))

```

```

: 0
: 8182

```

```

:
count      8182.000000
mean         3.416634
std          1.570037
min          0.100000
25%          2.300000
50%          3.700000
75%          4.500000
max          12.500000
Name: produkt, dtype: float64

```

```

:
count      8182.000000
mean         3.416634
std          1.570037
min          0.100000
25%          2.300000
50%          3.700000

```

```
75%          4.500000
max          12.500000
Name: produkt, dtype: float64
```

```
: 8182
```

```
:
: 0 (0.00%)
Sum: 0 (0.00%)
Ca: 0 (0.00%)
Mg: 0 (0.00%)
pH_KCL: 0 (0.00%)
P: 0 (0.00%)
K: 0 (0.00%)
N: 0 (0.00%)
Org: 0 (0.00%)
produkt: 18 (0.22%)
```

```
: 18
: 0.22%
: 8164
: 18
```

```
...
```

```
:
{'depth': 6, 'iterations': 500, 'learning_rate': 0.05}
```

```
CatBoost ( ):
RMSE: 0.7501
R2 Score: 0.7664
MAPE: 31.37%
```

```
:
      feature  importance
0  Agrotech    39.991141
2              25.191873
5         Mg    10.762833
3         Sum     6.534395
1         Soil     6.460281
7          P     2.692719
9          N     2.330991
8          K     2.027546
10         Org     1.801065
4          Ca     1.679669
6    pH_KCL     0.527487
```

```
[3]: import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
from sklearn.metrics import mean_squared_error, r2_score

def visualize_model_results(model, X_train, X_test, y_train, y_test, y_pred,
    features):
    #
    plt.figure(figsize=(20, 15))

    # 1.  vs
    plt.subplot(2, 2, 1)
    plt.scatter(y_test, y_pred, alpha=0.5)
    plt.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'r--',
    lw=2)
    plt.xlabel(' ')
    plt.ylabel(' ')
    plt.title(' ')

    # 2.
    plt.subplot(2, 2, 2)
    residuals = y_test - y_pred
    plt.scatter(y_pred, residuals, alpha=0.5)
    plt.axhline(y=0, color='r', linestyle='--')
    plt.xlabel(' ')
    plt.ylabel(' ')
    plt.title(' ')

    # 3.
    plt.subplot(2, 2, 3)
    plt.hist(residuals, bins=50)
    plt.xlabel(' ')
    plt.ylabel(' ')
    plt.title(' ')

    # 4.
    importance = pd.DataFrame({
        'feature': features,
        'importance': model.feature_importances_
    }).sort_values('importance', ascending=True)

    plt.subplot(2, 2, 4)
    plt.barh(importance['feature'], importance['importance'])
    plt.xlabel(' ')
    plt.title(' ')

    plt.tight_layout()
```

```

plt.show()

# 5.
plt.figure(figsize=(12, 8))
numeric_cols = X_train.select_dtypes(include=[np.number]).columns
correlation_matrix = X_train[numeric_cols].corr()
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', center=0)
plt.title(' ')
plt.tight_layout()
plt.show()

# 6. ( )
if hasattr(model, 'evals_result_'):
    plt.figure(figsize=(10, 6))
    # ,
    eval_metrics = model.evals_result_['learn'].keys()
    for metric in eval_metrics:
        plt.plot(model.evals_result_['learn'][metric],
                 label=f'{metric} (train)')
        if 'test' in model.evals_result_:
            plt.plot(model.evals_result_['test'][metric],
                     label=f'{metric} (test)')

    plt.xlabel('')
    plt.ylabel(' ')
    plt.title(' ')
    plt.legend()
    plt.show()

#
print("\n :")
for param_name, param_value in model.get_params().items():
    print(f"{param_name}: {param_value}")

#
print("\n :")
print(f" : {np.mean(residuals):.4f}")
print(f" : {np.std(residuals):.4f}")
print(f" : {np.min(residuals):.4f}")
print(f" : {np.max(residuals):.4f}")

#
print("\n :")
print(f" (MAE): {np.mean(np.abs(residuals)):.4f}")
print(f" : {np.median(np.abs(residuals)):.4f}")

#
percentiles = [1, 5, 25, 50, 75, 95, 99]

```



```

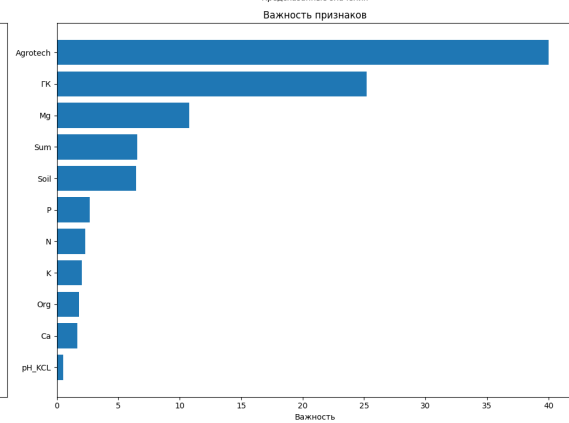


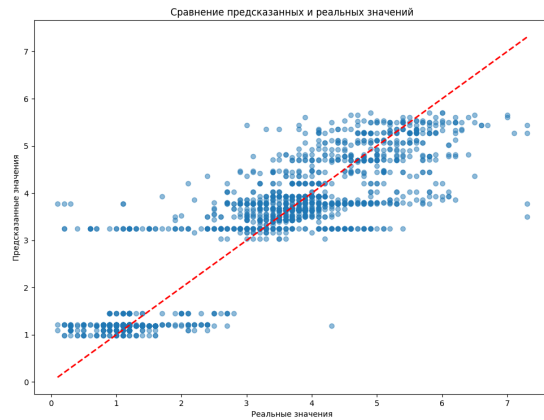
print("\n :")
for p in percentiles:
    print(f"{p}- : {np.percentile(residuals, p):.4f}")



#
visualize_model_results(
    model=best_model,
    X_train=X_train,
    X_test=X_test,
    y_train=y_train,
    y_test=y_test,
    y_pred=y_pred,
    features=features
)

#
print("\n :")
for cat_feature in ['Agrotech', 'Soil']:
    print(f"\n {cat_feature}:")
    for unique_val in X_test[cat_feature].unique():
        mask = X_test[cat_feature] == unique_val
        if mask.sum() > 0:
            group_rmse = np.sqrt(mean_squared_error(y_test[mask], y_pred[mask]))
            group_r2 = r2_score(y_test[mask], y_pred[mask])
            group_mape = calculate_mape(y_test[mask], y_pred[mask])
            print(f"\n {unique_val}:")
            print(f"RMSE: {group_rmse:.4f}")
            print(f"R2: {group_r2:.4f}")
            print(f"MAPE: {group_mape:.2f}%")
            print(f" : {mask.sum()}")

#
best_model.save_model('catboost_model.cbm')
print("\n 'catboost_model.cbm'")

```





```
:
iterations: 500
learning_rate: 0.05
depth: 6
loss_function: RMSE
random_seed: 42
verbose: False
cat_features: ['Agrotech', 'Soil']
```

```
:
: -0.0105
: 0.7500
: -3.6772
: 3.8020
```

```
:
(MAE): 0.5189
: 0.3333
```

```
:
1- : -2.2523
5- : -1.1702
25- : -0.3578
50- : 0.0010
75- : 0.3150
95- : 1.2477
99- : 1.8900
```

```
:
```

Agrotech:

```
X_1:
RMSE: 0.7468
R2: 0.0723
MAPE: 14.21%
: 94
```

```
K:
RMSE: 0.4671
R2: 0.8653
MAPE: 20.56%
: 408
```

TZ:

RMSE: 0.9659  
R2: 0.4082  
MAPE: 41.79%  
: 617

VI:  
RMSE: 0.6263  
R2: 0.8678  
MAPE: 30.58%  
: 514

Soil:

:  
RMSE: 0.7493  
R2: 0.7808  
MAPE: 40.37%  
: 435

:  
RMSE: 0.5874  
R2: 0.8521  
MAPE: 23.46%  
: 633

:  
RMSE: 0.7089  
R2: 0.5012  
MAPE: 12.57%  
: 218

:  
RMSE: 0.5419  
R2: 0.7133  
MAPE: 10.49%  
: 26

:  
RMSE: 1.1338  
R2: 0.1972  
MAPE: 64.17%  
: 225

:  
RMSE: 0.7577  
R2: 0.1625  
MAPE: 14.76%  
: 70

:  
RMSE: 0.5459  
R2: 0.0437  
MAPE: 14.19%  
: 13

Unknown:  
RMSE: 0.5918  
R2: -2.6654  
MAPE: 10.32%  
: 3

:  
RMSE: 0.6738  
R2: 0.1217  
MAPE: 11.99%  
: 10

'catboost\_model.cbm'

```
[4]: import os
# Graphviz
os.environ["PATH"] += os.pathsep + r"C:\Program Files\Graphviz\bin"

def visualize_trees(model, features, max_trees=3):
    """
        CatBoost
    """
    try:
        import graphviz

        # dot
        try:
            from subprocess import run, PIPE
            result = run(['dot', '-V'], stdout=PIPE, stderr=PIPE)
            print("Graphviz version:", result.stderr.decode())
        except Exception as e:
            print("Graphviz:", str(e))

        print(f"\n {max_trees} ")

        #
        if not os.path.exists('tree_visualizations'):
            os.makedirs('tree_visualizations')

        # dot
```

```

graphviz.backend.executables = {'dot': r'C:\Program_
↪Files\Graphviz\bin\dot.exe'}

#
for tree_idx in range(min(max_trees, model.tree_count_)):
    try:
        tree_graph = model.plot_tree(tree_idx)

        # dot
        output_file = f'tree_visualizations/tree_{tree_idx}'
        tree_graph.render(filename=output_file,
                           format='png',
                           cleanup=True,
                           engine='dot')
        print(f" {tree_idx} {output_file}.png")

    except Exception as e:
        print(f" {tree_idx}: {str(e)}")

# ...

except Exception as e:
    print(f" : {str(e)}")

#
visualize_trees(best_model, features, max_trees=3)

```

Graphviz version: dot - graphviz version 12.2.1 (20241206.2353)

```

3
0  tree_visualizations/tree_0.png
1  tree_visualizations/tree_1.png
2  tree_visualizations/tree_2.png

```

```

[5]: from sklearn.ensemble import RandomForestRegressor, GradientBoostingRegressor
from xgboost import XGBRegressor
from lightgbm import LGBMRegressor
from sklearn.metrics import mean_squared_error, r2_score
import pandas as pd
import numpy as np

#
results = {}

# 1. CatBoost
catboost_model = CatBoostRegressor(

```

```

        cat_features=cat_features,
        iterations=2000,
        depth=8,
        learning_rate=0.03,
        verbose=False,
        random_seed=42
    )

# 2. XGBoost
xgb_model = XGBRegressor(
    n_estimators=2000,
    max_depth=8,
    learning_rate=0.03,
    random_state=42,
    n_jobs=-1
)

# 3. LightGBM
lgb_model = LGBMRegressor(
    n_estimators=2000,
    max_depth=8,
    learning_rate=0.03,
    random_state=42,
    n_jobs=-1
)

# 4. Random Forest
rf_model = RandomForestRegressor(
    n_estimators=500,
    max_depth=8,
    random_state=42,
    n_jobs=-1
)

# 5. Gradient Boosting
gb_model = GradientBoostingRegressor(
    n_estimators=500,
    max_depth=8,
    learning_rate=0.03,
    random_state=42
)

#
models = {
    'CatBoost': catboost_model,
    'XGBoost': xgb_model,
    'LightGBM': lgb_model,

```



```

    'Random Forest': rf_model,
    'Gradient Boosting': gb_model
}

#         CatBoost
from sklearn.preprocessing import LabelEncoder

#
X_train_encoded = X_train.copy()
X_test_encoded = X_test.copy()

#
label_encoders = {}
for feature in cat_features:
    le = LabelEncoder()
    X_train_encoded[feature] = le.fit_transform(X_train[feature])
    X_test_encoded[feature] = le.transform(X_test[feature])
    label_encoders[feature] = le

#
for name, model in models.items():
    print(f"\n {name}...")

    #
    if name == 'CatBoost':
        model.fit(X_train, y_train)
        y_pred = model.predict(X_test)
    else:
        model.fit(X_train_encoded, y_train)
        y_pred = model.predict(X_test_encoded)

    #
    rmse = np.sqrt(mean_squared_error(y_test, y_pred))
    r2 = r2_score(y_test, y_pred)
    mape = calculate_mape(y_test, y_pred)

    results[name] = {
        'RMSE': rmse,
        'R2': r2,
        'MAPE': mape
    }

    print(f"{name} :")
    print(f"RMSE: {rmse:.4f}")
    print(f"R2: {r2:.4f}")
    print(f"MAPE: {mape:.2f}%")

```

```

#
import matplotlib.pyplot as plt

# RMSE
plt.figure(figsize=(12, 6))
plt.bar([name for name in results.keys()],
        [results[name]['RMSE'] for name in results.keys()])
plt.title(' RMSE ')
plt.xticks(rotation=45)
plt.ylabel('RMSE')
plt.tight_layout()
plt.show()

# R2
plt.figure(figsize=(12, 6))
plt.bar([name for name in results.keys()],
        [results[name]['R2'] for name in results.keys()])
plt.title(' R2 ')
plt.xticks(rotation=45)
plt.ylabel('R2')
plt.tight_layout()
plt.show()

# DataFrame
results_df = pd.DataFrame.from_dict(results, orient='index')
print("\n :")
print(results_df)

#
results_df.to_csv('model_comparison_results.csv')

```

```

CatBoost...
CatBoost :
RMSE: 0.7535
R2: 0.7642
MAPE: 31.31%

```

```

XGBoost...
XGBoost :
RMSE: 0.7497
R2: 0.7666
MAPE: 31.19%

```

```

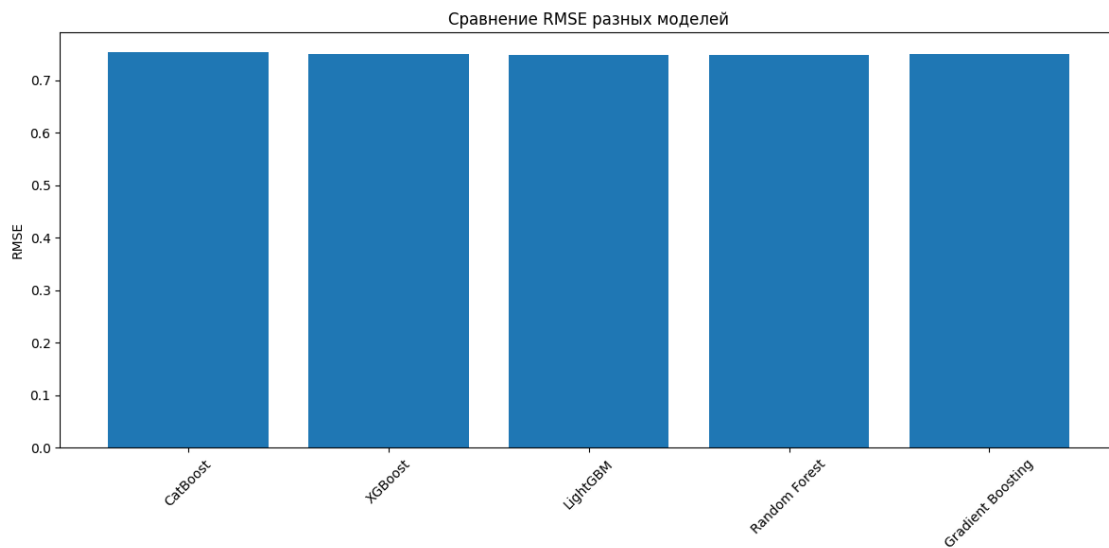
LightGBM...
[LightGBM] [Info] Auto-choosing row-wise multi-threading, the overhead of
testing was 0.000401 seconds.
You can set `force_row_wise=true` to remove the overhead.

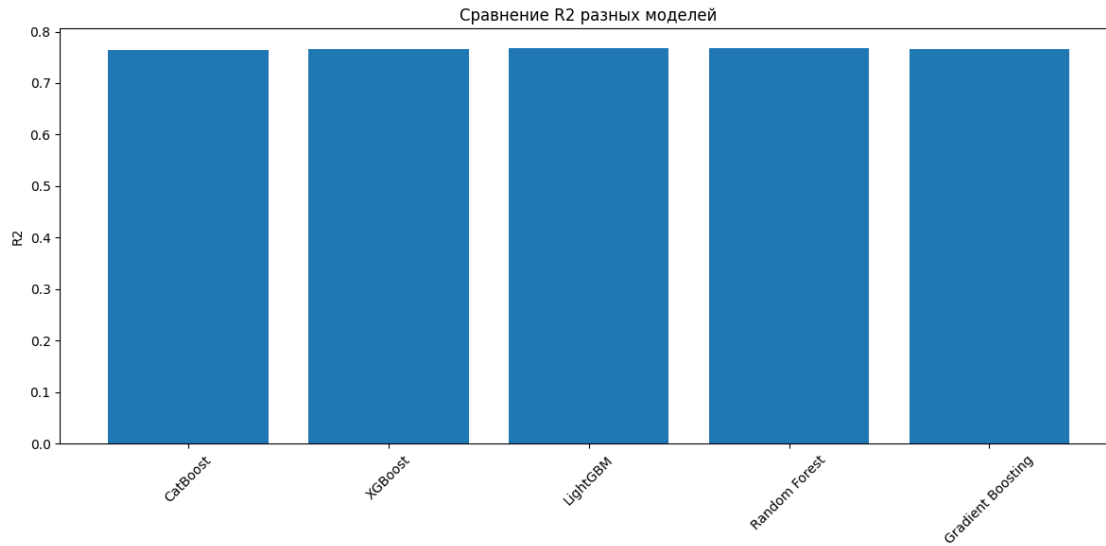
```

And if memory is not enough, you can set `force\_col\_wise=true`.  
[LightGBM] [Info] Total Bins 82  
[LightGBM] [Info] Number of data points in the train set: 6531, number of used features: 11  
LightGBM :  
RMSE: 0.7483  
R2: 0.7674  
MAPE: 31.18%

Random Forest...  
Random Forest :  
RMSE: 0.7491  
R2: 0.7670  
MAPE: 31.21%

Gradient Boosting...  
Gradient Boosting :  
RMSE: 0.7502  
R2: 0.7663  
MAPE: 31.20%





```
:
```

	RMSE	R2	MAPE
CatBoost	0.753495	0.764226	31.305141
XGBoost	0.749670	0.766614	31.186627
LightGBM	0.748340	0.767441	31.177163
Random Forest	0.749088	0.766976	31.213115
Gradient Boosting	0.750208	0.766278	31.200033

```
[6]: #
tree_graph = best_model.plot_tree(
    tree_index=0,
    pool=None
)
tree_graph.render(filename='first_tree',
                   format='png',
                   cleanup=True)
```

```
-----
TypeError                                Traceback (most recent call last)
Cell In[6], line 2
      1 #
----> 2 tree_graph = best_model.plot_tree(
      3     tree_index=0,
      4     pool=None
      5 )
      6 tree_graph.render(filename='first_tree',
      7                     format='png',
      8                     cleanup=True)
```

```
TypeError: CatBoost.plot_tree() got an unexpected keyword argument 'tree_index'
```

```
[ ]:
```